

Training a Real Robot Arm

To Scoop Some Beans

Isaac Hirsch, Matthew LeBar, and Chad Schmerling

The Project Idea

Goal:

- Use a low-cost robot arm to dig a hole in a tub filled with uncooked beans
- Make use of a leader arm to train a digging policy using imitation learning
- Use a GoPro camera to provide the policy with feedback

Considered alternative goals: teaching arm to disconnect itself, teaching arm to twist fingers, etc.

Resources:

- Koch robot follower arm
- Koch robot leader arm
- GoPro Hero Black 7
- Plastic bin
- Uncooked pinto beans

Environment Setup

- Our environment was set up in the TTIC DuckieTown Robotics Lab (courtesy of Professor Matthew “Matt” Walter)
- We clamped the follower arm down next to the bin, which was velcroed to the table to keep it stable
- The leader arm was clamped down nearby for easy observation. We 3d printed a support to set it at a similar height
- We set up the GoPro right next to the bin to provide visual information for the model, and marked it with tape so we could keep its position consistent




Data collection

- To train and evaluate our arm's performance, we used Hugging Face's LeRobot repo, which provided calibration, training, and testing scripts
- Our idea was to start with imitation learning to provide a foundation, then perform reinforcement learning from there
- However, just completing imitation learning was quite challenging
- We used the leader arm to teleoperate the follower arm for 29 episodes, each 20 seconds long
- The motor overheated and we struggled to find a way to combine multiple series of episodes
- The operator used the camera feed to try and scoop out as many beans from the center of the bin as possible



Model

- We used the *Action Chunking with Transformers* (ACT) algorithm built into the LeRobot repository for imitation learning
- ACT uses a transformer based VAE to learn a policy



Training

- AdamW
 - lr = 1e-5
 - batchsize = 16
 - l2 weight decay = 0.0001
- Gradient Clipping
- Checkpointed every 100 steps
- Evaluated IRL every 200 steps
- Trained for 13500 steps

Algorithm 1 ACT Training

- 1: Given: Demo dataset \mathcal{D} , chunk size k , weight β .
 - 2: Let a_t, o_t represent action and observation at timestep t , \bar{o}_t represent o_t without image observations.
 - 3: Initialize encoder $q_\phi(z|a_{t:t+k}, \bar{o}_t)$
 - 4: Initialize decoder $\pi_\theta(\hat{a}_{t:t+k}|o_t, z)$
 - 5: **for** iteration $n = 1, 2, \dots$ **do**
 - 6: Sample $o_t, a_{t:t+k}$ from \mathcal{D}
 - 7: Sample z from $q_\phi(z|a_{t:t+k}, \bar{o}_t)$
 - 8: Predict $\hat{a}_{t:t+k}$ from $\pi_\theta(\hat{a}_{t:t+k}|o_t, z)$
 - 9: $\mathcal{L}_{reconst} = MSE(\hat{a}_{t:t+k}, a_{t:t+k})$
 - 10: $\mathcal{L}_{reg} = D_{KL}(q_\phi(z|a_{t:t+k}, \bar{o}_t) \parallel \mathcal{N}(0, I))$
 - 11: Update θ, ϕ with ADAM and $\mathcal{L} = \mathcal{L}_{reconst} + \beta\mathcal{L}_{reg}$
-

Obstacles and Solutions

- ARM mechanical issues:
 - replacing a circuit board
 - needing to frequently power cycle robot arms
 - 3D Printing replacement parts
- GoPro required multiple specialized connectors and a backdoor exploit to work as a webcam
- LeRobot repo is extremely intricate, tailoring it to work for our particular project was often quite involved

Evaluation 1


- We started off with a dumbed-down ACT model to expedite training.
 - There were *still* 11M parameters to train
- It was unable to learn






Evaluation 2


- We then returned to the author's sets
 - There were now 45M parameters to train
- It learned much better
















Goal



Ours (13,500)



Next steps

- Improvements for shuffling
 - Train for longer to get better results for imitation learning
 - Obtain more samples of imitation learning
 - Use reinforcement learning to bolster shuffling technique
- Future Directions
 - Use imitation learning as a base to start reinforcement learning